


Introduction to Machine Learning

CART: In a nutshell



Learning goals

- Understand basic structure of CART models
- Understand basic concepts used to fit CART models

WHAT IS A TREE?

Basic idea:

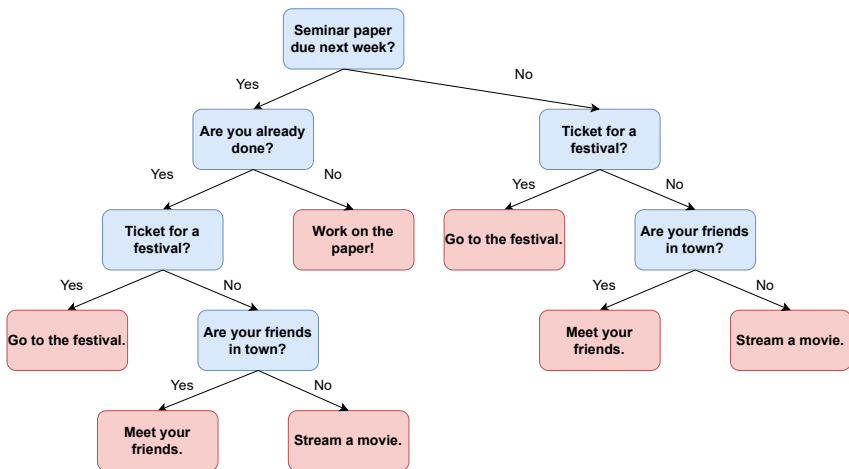
- Divide feature space into sub-regions.
- For each region, learn best constant prediction from training data.

Classification **A**nd **R**egression **T**rees are a class of models that can:

- model non-linear feature effects
- facilitate interactions of features
- be inherently interpreted

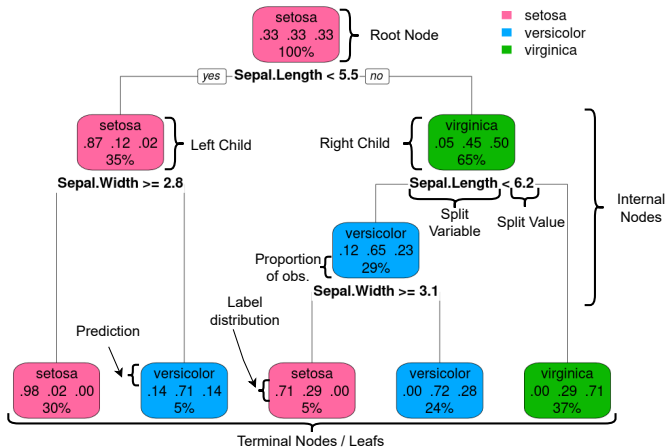
WHAT IS A TREE?

A decision tree is a set of hierarchical binary partitions, e.g., your evening planning decision (target) could be based on a decision tree:



WHAT IS A TREE?


Instead of life choices, we can predict the Species of flowers described in the iris data set using features `Sepal.Width` and `Sepal.Length`.



CART AS A PREDICTOR

Instead of the visual description, we can also describe trees through their division of the feature space \mathcal{X} into **rectangular regions**, Q_m :


$$f(\mathbf{x}) = \sum_{m=1}^M c_m \mathbb{I}(\mathbf{x} \in Q_m),$$





TASKS FOR CART

- CARTs can have categorical and numerical targets.
- In both cases, the leafs, i.e., the ultimate nodes, define the predictions.

Categorical target:




Numerical target:



HOW TO FIT A CART


- A *recursive* greedy search in the feature space optimizes CARTs
- In each iteration, the best split is selected

Iteration 1:




HOW TO FIT A CART

Iteration 2:




Iteration 3:



HOW TO FIT A CART

- This procedure can run until each observation has its own leaf.
- Then, the tree will not generalize well and overfit:



- Thus, we need techniques to keep the tree small and informative.
- ⇒ In practice, trees are often used as base learners for ensemble learners like Random Forests.