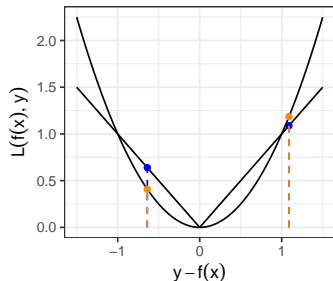


Introduction to Machine Learning

Supervised Regression: Linear Models with L_1 Loss

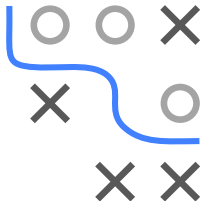
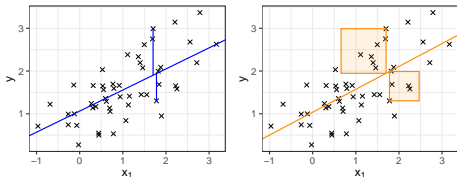


Learning goals

- Understand difference between L_1 and L_2 regression
- See how choice of loss affects optimization & robustness

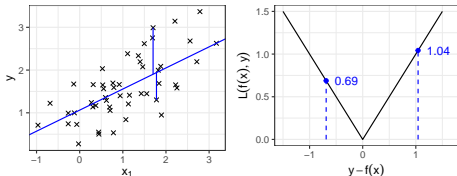
ABSOLUTE LOSS

- L_2 regression minimizes quadratic residuals – wouldn't **absolute** residuals seem more natural?

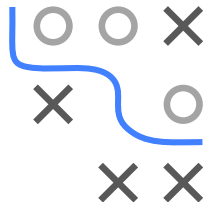




- L_1 loss / absolute error / least absolute deviation (LAD)

$$L(y, f(\mathbf{x})) = |y - f(\mathbf{x})|$$



L1 VS L2 – LOSS SURFACE



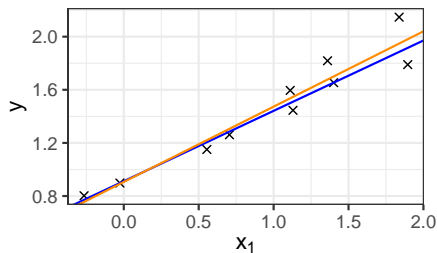
L1 loss (left) harder to optimize than L2 loss (right)

- Convex but **not differentiable** in $y - f(\mathbf{x}) = 0$
- No analytical solution

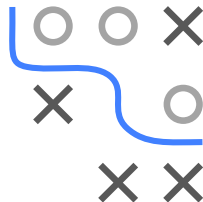
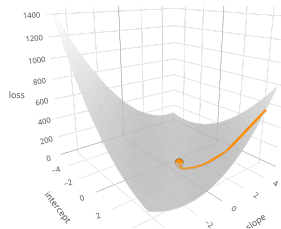
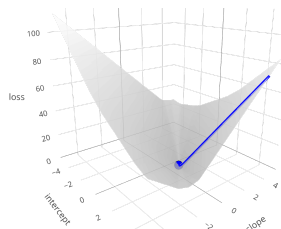
L1 VS L2 – ESTIMATED PARAMETERS

- Results of $L1$ and $L2$ regression often not that different
- Simulated data: $y^{(i)} = 1 + 0.5x_1^{(i)} + \epsilon^{(i)}$, $\epsilon^{(i)} \stackrel{i.i.d}{\sim} \mathcal{N}(0, 0.01)$

| | intercept | slope |
|------|-----------|-------|
| $L1$ | 0.91 | 0.53 |
| $L2$ | 0.91 | 0.57 |





absolute quadratic




L1 VS L2 – ROBUSTNESS

- L2 quadratic in residuals \rightsquigarrow outlying points carry lots of weight
- E.g., $3\times$ residual $\Rightarrow 9\times$ loss contribution
- L1 more **robust** in presence of outliers (example ctd.):




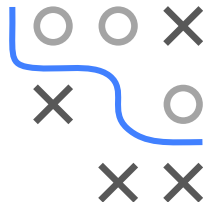
absolute quadratic



absolute quadratic

L1 VS L2 – OPTIMIZATION COST

- Real-world weather problem \rightsquigarrow predict mean temperature
- Compare **time** to fit L1 (`quantreg::rq()`) vs L2 (`lm::lm()`) for different dataset proportions (repeat 50 \times)



Loss

| | Fitted: <i>L1</i> | Fitted: <i>L2</i> |
|----------------------|--------------------|--------------------|
| Total <i>L1</i> loss | 8.98×10^4 | 8.99×10^4 |
| Total <i>L2</i> loss | 5.83×10^6 | 5.81×10^6 |

Estimated coefficients

| x_j | <i>L1</i> : $\hat{\theta}_j$ | <i>L2</i> : $\hat{\theta}_j$ |
|-----------------|------------------------------|------------------------------|
| Max_temperature | 0.553 | 0.563 |
| Min_temperature | 0.441 | 0.427 |
| Visibility | 0.026 | 0.041 |
| Wind_speed | 0.002 | 0.010 |
| Max_wind_speed | -0.026 | -0.039 |
| (Intercept) | -0.380 | -0.102 |

L1 slower to optimize!