

Exercise 1: SVM – Regression

For the data set $\mathcal{D} = ((\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)}))$ with $y^{(i)} \in \mathbb{R}$, assume that for a fixed $\epsilon > 0$ all observations are within the ϵ -tube around $f(\mathbf{x} \mid \theta_0, \boldsymbol{\theta}) = \boldsymbol{\theta}^\top \mathbf{x} + \theta_0$ for any $(\theta_0, \boldsymbol{\theta})^\top \in \tilde{\Theta}$, i.e.,

$$y^{(i)} \in [f(\mathbf{x}^{(i)} \mid \theta_0, \boldsymbol{\theta}) - \epsilon, f(\mathbf{x}^{(i)} \mid \theta_0, \boldsymbol{\theta}) + \epsilon], \quad \forall i \in \{1, \dots, n\}, \quad \forall (\theta_0, \boldsymbol{\theta})^\top \in \tilde{\Theta},$$

where $\tilde{\Theta} \subset \mathbb{R}^{p+1}$ is some non-empty parameter subset. Let

$$d_\epsilon \left(f(\cdot \mid \theta_0, \boldsymbol{\theta}), \mathbf{x}^{(i)} \right) := \epsilon - |y^{(i)} - f(\mathbf{x}^{(i)} \mid \theta_0, \boldsymbol{\theta})| = \epsilon - |y^{(i)} - \boldsymbol{\theta}^\top \mathbf{x}^{(i)} - \theta_0|$$

be the (signed) ϵ -distance of the prediction error. The maximal ϵ -distance of the prediction error of f to the whole data set \mathcal{D} is

$$\gamma_\epsilon = \max_{i=1, \dots, n} \left\{ d_\epsilon \left(f(\cdot \mid \theta_0, \boldsymbol{\theta}), \mathbf{x}^{(i)} \right) \right\}.$$

(a) Let $(\theta_0, \boldsymbol{\theta})^\top \in \mathbb{R}^{p+1}$ be arbitrary. Which (type of) values does $d_\epsilon \left(f(\cdot \mid \theta_0, \boldsymbol{\theta}), \mathbf{x}^{(i)} \right)$ have if $y^{(i)}$ is

- within the ϵ -tube around $f(\mathbf{x}^{(i)} \mid \theta_0, \boldsymbol{\theta})$?
- not within the ϵ -tube around $f(\mathbf{x}^{(i)} \mid \theta_0, \boldsymbol{\theta})$?

What would be a desirable choice of the parameters $(\theta_0, \boldsymbol{\theta})^\top$ with respect to γ_ϵ ? Is the choice of the parameters unique in general?

(b) Argue that

$$\begin{aligned} \min_{\theta_0, \boldsymbol{\theta}} \quad & \frac{1}{2} \|\boldsymbol{\theta}\|^2 \\ \text{s.t.} \quad & \epsilon - y^{(i)} + \boldsymbol{\theta}^\top \mathbf{x}^{(i)} + \theta_0 \geq 0 \quad \forall i \in \{1, \dots, n\} \\ \text{and} \quad & \epsilon + y^{(i)} - \boldsymbol{\theta}^\top \mathbf{x}^{(i)} - \theta_0 \geq 0 \quad \forall i \in \{1, \dots, n\}. \end{aligned}$$

is a suitable optimization problem for the desired choice in (a).

- (c) Derive the Lagrange function $L(\boldsymbol{\theta}, \theta_0, \boldsymbol{\alpha})$ of the optimization problem as well as its dual form.
- (d) Find the stationary points of L . What can be inferred from the solution of the dual problem?
- (e) Derive the “softened” version of the optimization problem in (b).
- (f) Rewrite the “softened” version of the optimization problem in (b) as a regularized empirical risk minimization problem for a suitable loss function for regression.

Exercise 2: SVM – Optimization

Write your own stochastic subgradient descent routine to solve the soft-margin SVM in the primal formulation.

Hints:

- Use the regularized-empirical-risk-minimization formulation, i.e., an optimization criterion without constraints.
- No kernels, just a linear SVM.
- Compare your implementation with an existing implementation (e.g., `kernlab` in R). Are your results similar? Note that you might have to switch off the automatic data scaling in the already existing implementation.